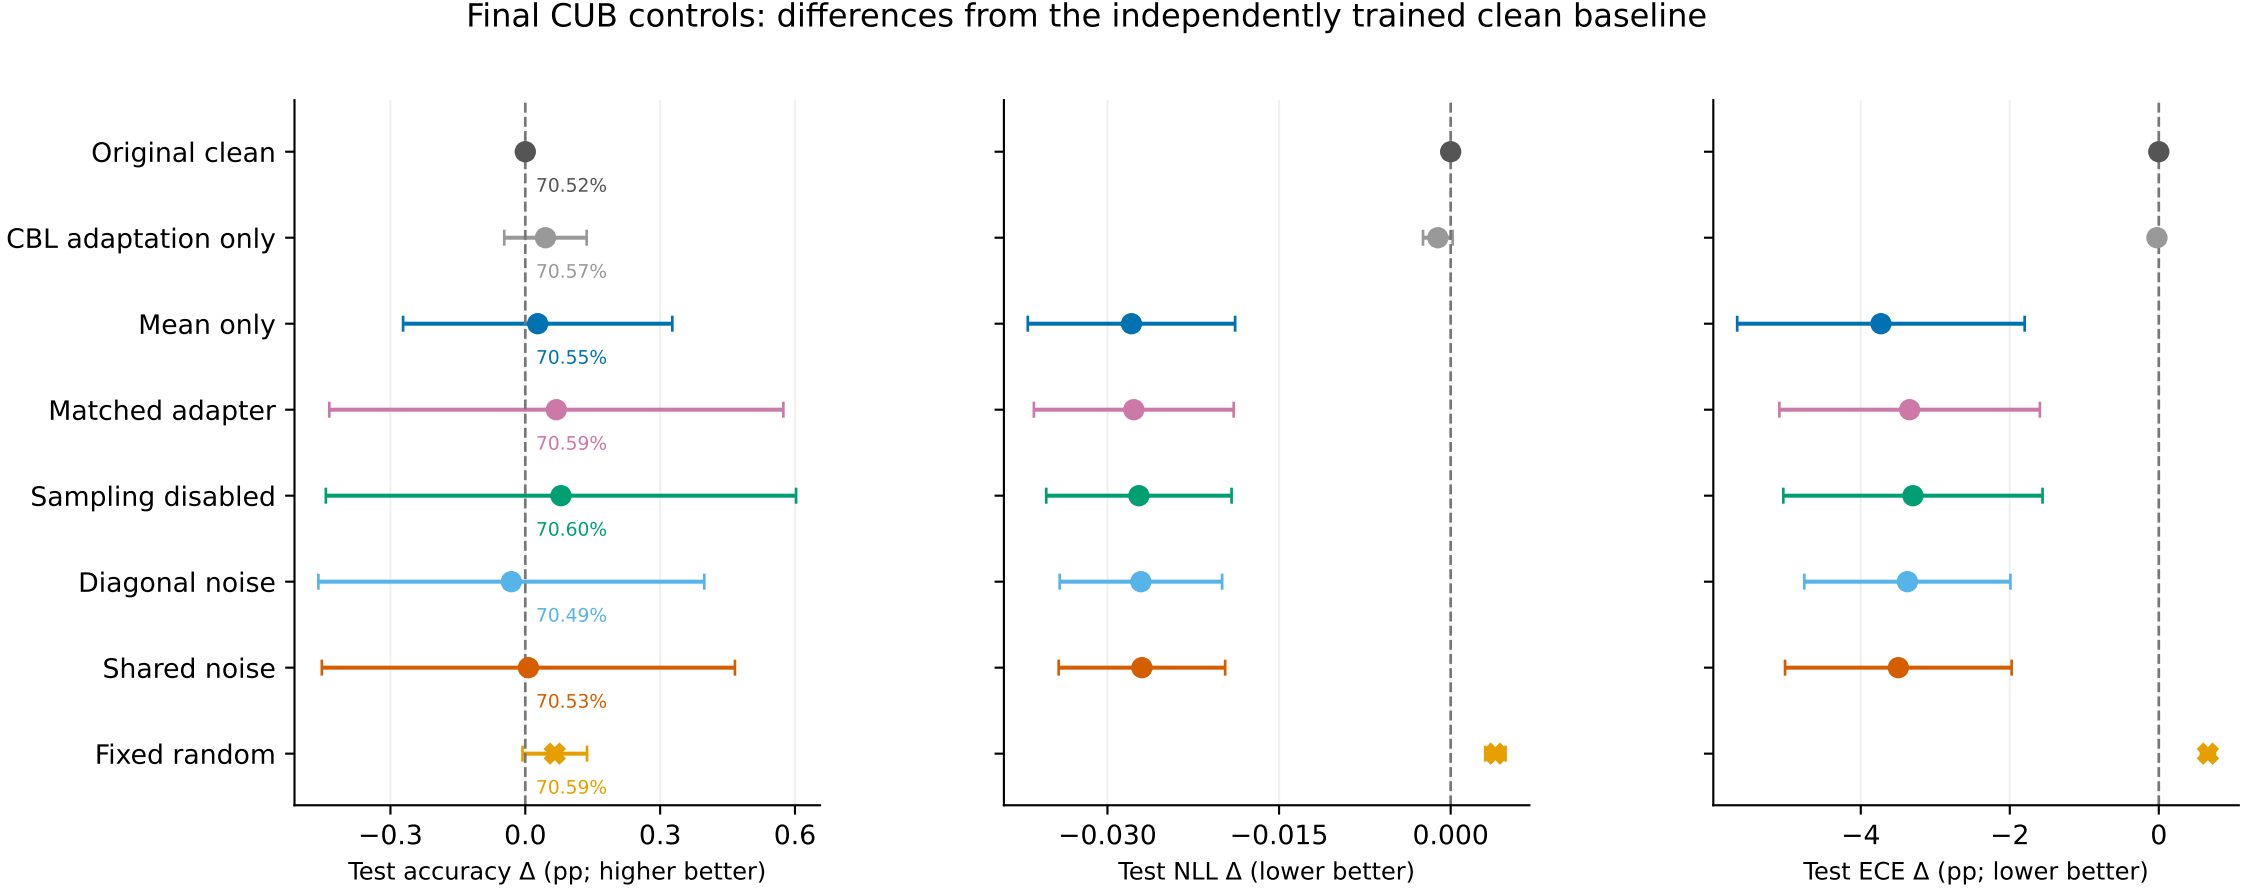
Final CUB controls: differences from the independently trained clean baseline
Original clean
CBL adaptation only
Mean only
Matched adapter
Sampling disabled
Diagonal noise
Shared noise
Fixed random
70.52%
70.57%
70.55%
70.59%
70.60%
70.49%
70.53%
70.59%
−0.3
0.0
0.3
0.6
Test accuracy Δ (pp; higher better)
−0.030
−0.015
0.000
Test NLL Δ (lower better)
−4
−2
0
Test ECE Δ (pp; lower better)
95% paired t intervals across five training seeds. X: at least one validation gate failed; retained as diagnostic control. No test selection.