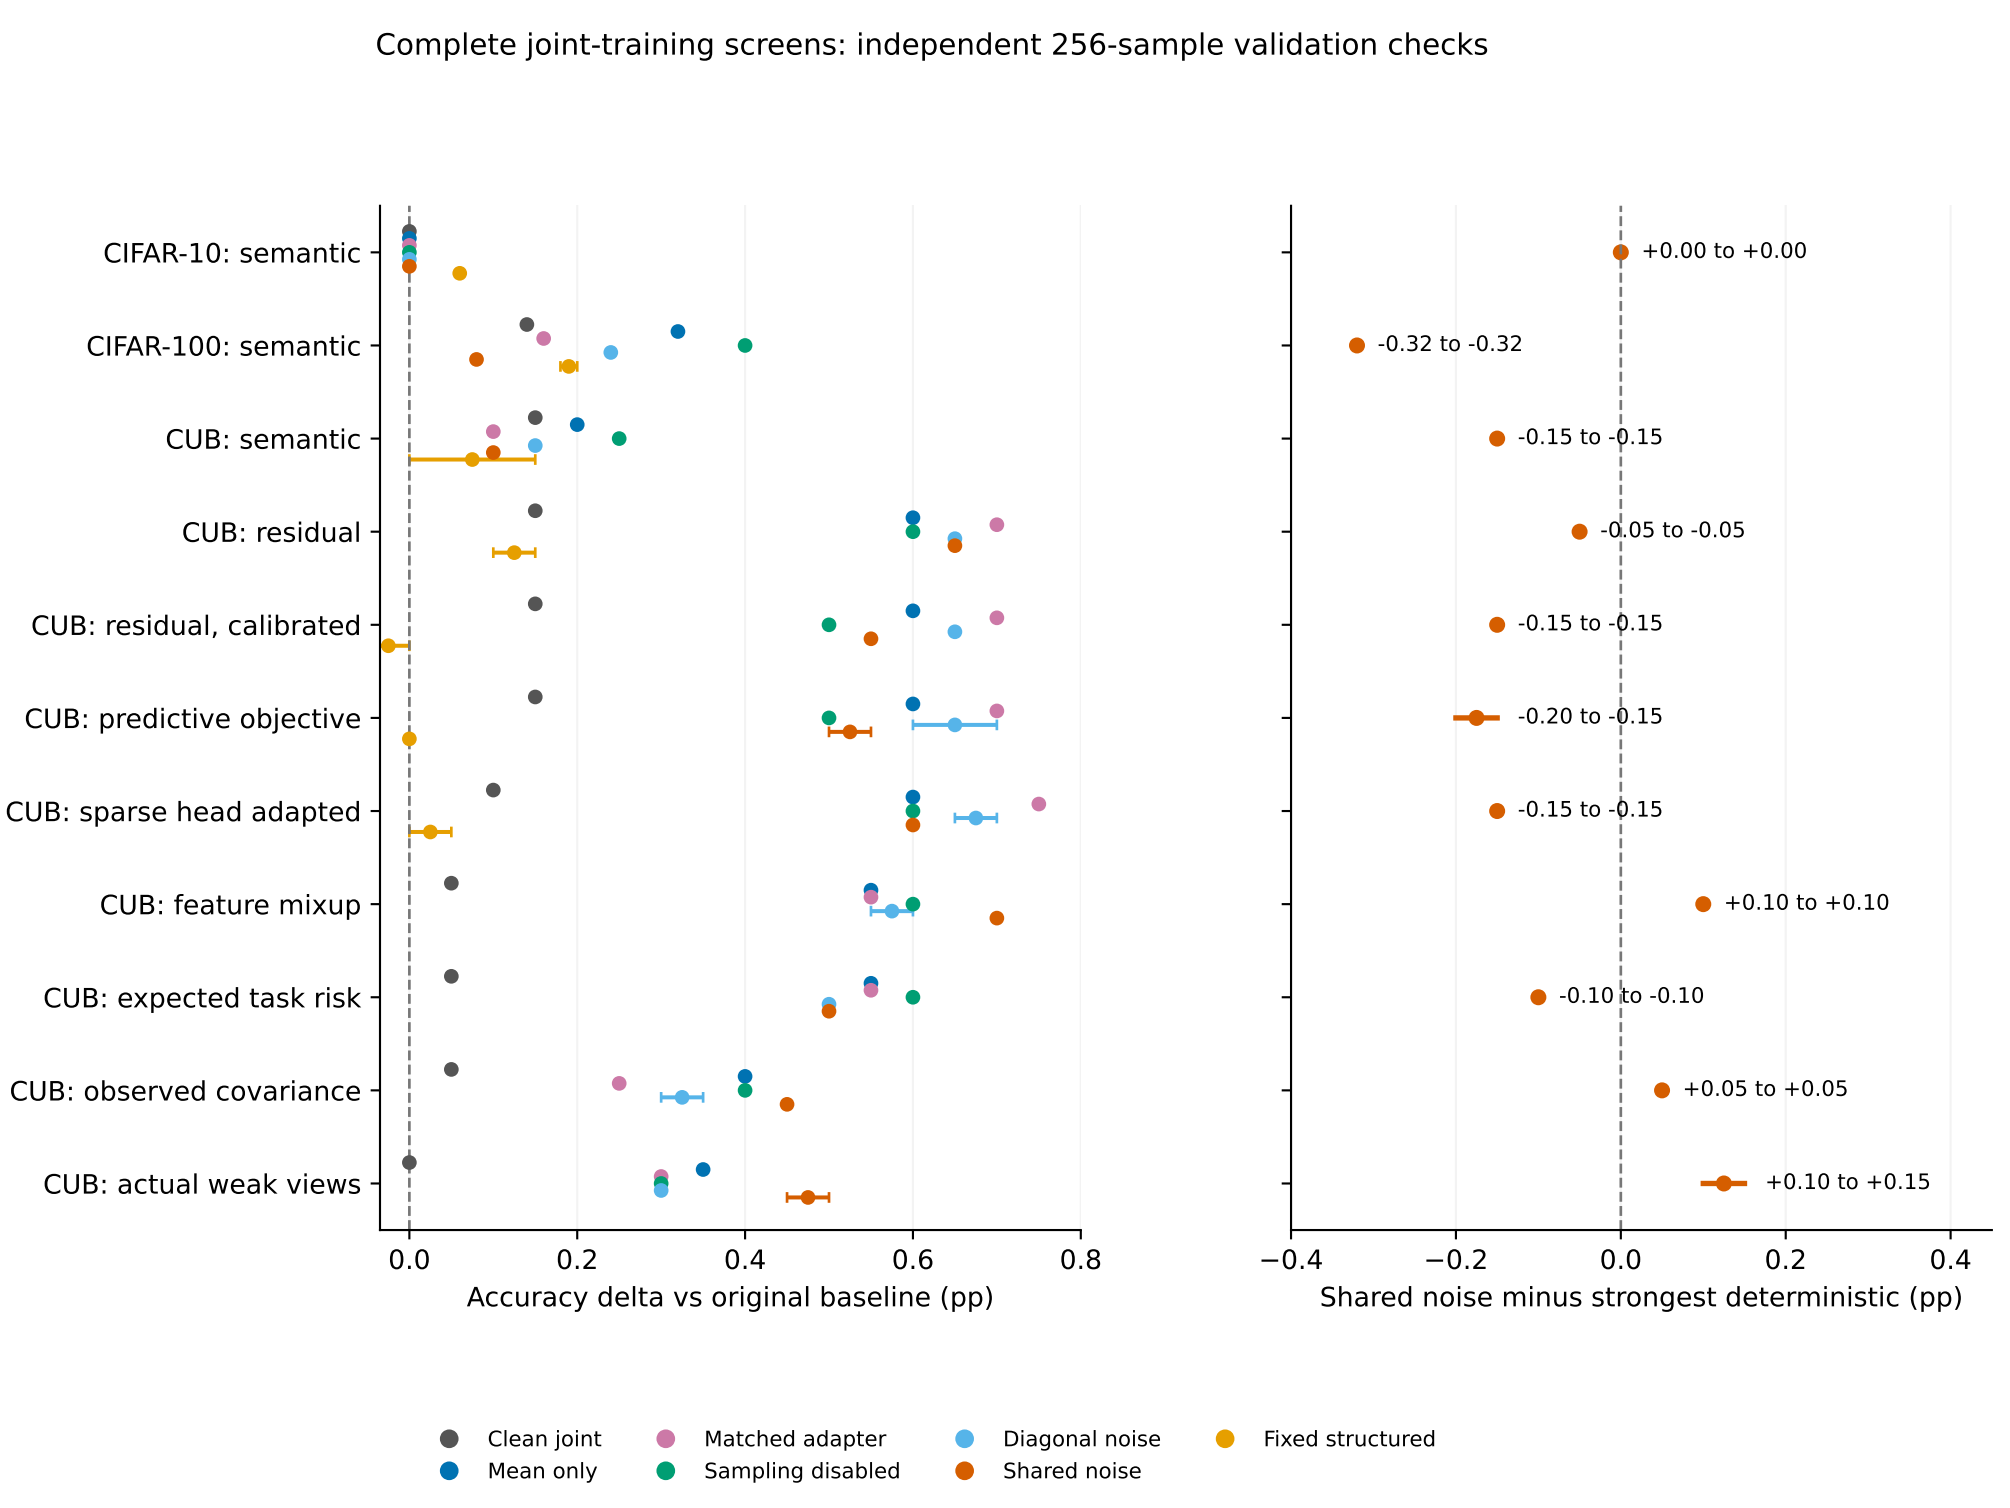
Complete joint-training screens: independent 256-sample validation checks
CIFAR-10: semantic
CIFAR-100: semantic
CUB: semantic
CUB: residual
CUB: residual, calibrated
CUB: predictive objective
CUB: sparse head adapted
CUB: feature mixup
CUB: expected task risk
CUB: observed covariance
CUB: actual weak views
0.0
0.2
0.4
0.6
0.8
Accuracy delta vs original baseline (pp)
+0.00 to +0.00
-0.32 to -0.32
-0.15 to -0.15
-0.05 to -0.05
-0.15 to -0.15
-0.20 to -0.15
-0.15 to -0.15
+0.10 to +0.10
-0.10 to -0.10
+0.05 to +0.05
+0.10 to +0.15
−0.4
−0.2
0.0
0.2
0.4
Shared noise minus strongest deterministic (pp)
Clean joint
Mean only
Matched adapter
Sampling disabled
Diagonal noise
Shared noise
Fixed structured
Cases selected by original validation MC; independent resampling never reselects epochs. Ranges are MC variation, not confidence intervals.